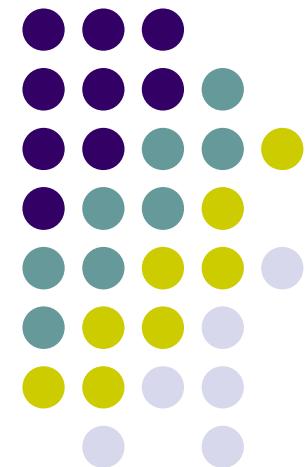
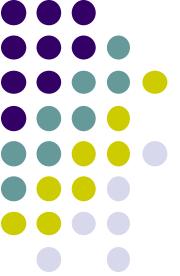


Computer Architecture

Assoc. Prof. Nguyễn Trí Thành, PhD
UNIVERSITY OF ENGINEERING AND TECHNOLOGY
FACULTY OF INFORMATION TECHNOLOGY
DEPARTMENT OF INFORMATION SYSTEMS
ntthanh@vnu.edu.vn



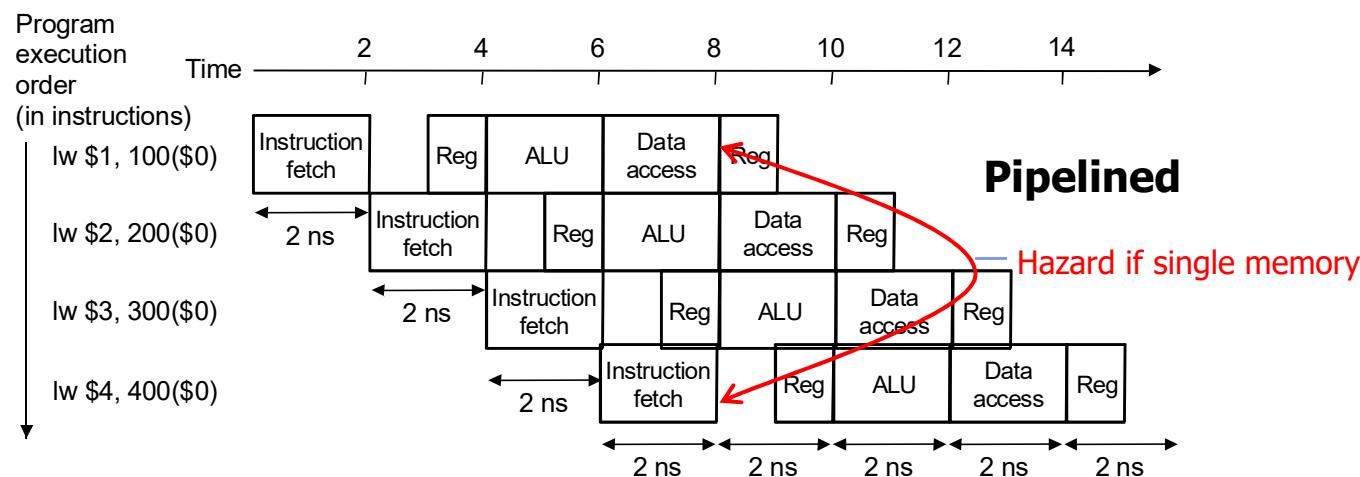


Chapter 4: Cache

Structural Hazards - Reminded



- *Structural hazard*: inadequate hardware to simultaneously support all instructions in the pipeline in the same clock cycle
- E.g., suppose *single – not separate* – instruction and data memory in pipeline below with *one read port*
 - then a structural hazard between first and fourth l_w instructions



- MIPS was designed to be pipelined: structural hazards are easy to avoid!

LENOVO B470 TECHNICAL SPECIFICATIONS

Processor	2 nd Generation Intel® Core™ i5-2410M Processor with Intel Turbo Boost Technology* 2.0 (2.30GHz, 3MB Cache) 2 nd Generation Intel® Core™ i3-2310M Processor (2.10GHz, 3MB Cache)
Operating System	Genuine Windows® 7 Professional Genuine Windows® 7 Home Premium Genuine Windows® 7 Home Basic
Display/Resolution	14.0" high-definition (1366*768), 16:9 widescreen
Video Graphics	Up to NVIDIA® GeForce® 410M 1GB graphics
Memory	Up to 8GB DDR3 1066/1333MHz memory
Hard Disk Drive	250GB/320GB/500GB/640GB/750GB/1TB 5400pm or 320GB/500GB/750GB/1TB 7200rpm HDD
Optical Drive	Integrated DVD reader/writer or Blu-ray Disc™ ¹ drive
Sound	2.0 speakers (2*1.5W)
Integrated Communications	802.11bg/bgn WiFi, 10/100M LAN, Bluetooth® (optional)

LENOV**Processor / Chipset**

	CPU	Intel 2nd Gen Core i5 2 i5-2410M / 2.3 GHz
Processor	Max Turbo Speed	2.9 GHz
Operating System	Number of Cores	Dual-Core
Display/Resolution	Cache	L3 cache - 3.0 MB
Video Graphics	64-bit Computing	Yes
Memory		
Hard Disk Drive	Features	Intel Turbo Boost Technology 2.0, Hyper-Threading Technology, Integrated memory controller
Optical Drive		
Sound		
Integrated Communications	Memory	
	RAM	4.0 GB (2 x 2 GB)
	Max RAM Supported	8.0 GB
	Technology	DDR3 SDRAM
	Speed	1333.0 MHz / PC3-10600
	Slots Qty	2

LENOV**Processor / Chipset**

CPU

Intel 2nd Gen Core i5 2 i5-2410M / 2.3 GHz

Processor**Max Turbo Speed**

2.9 GHz

Op

Now Here are some Tech Specs for my fellow nerds:

Dis

CPU: i7 2670

Vic

4 core 8 Thread

Me

2.2ghz 3.1ghz turbo

Ha

32nm 45watt TDP

Op

5 GT/s x22 multiplier

So

Socket G2 (rPGA988B)

Int**Caches:****Co**

L1 data: 128kb

echnology 2.0,
echnology,
controller

L1 inst: 128kb

L2: 1024kb

L3: 6mb

Bus speed: 99.8mhz

MAX RAM SUPPORTED

0.0 GB

Technology

DDR3 SDRAM

Speed

1333.0 MHz / PC3-10600

Slots Qty

2

	Processor / Chipset	Architecture / Microarchitecture
Processor	Microarchitecture	Sandy Bridge
	Processor core ?	Sandy Bridge
	Core stepping ?	D2 (Q1HK, SR00B)
	Manufacturing process	0.032 micron High-K metal gate process
Op	Now H	216mm ²
Op	Data width	64 bit
Dis	CPU: i7	The number of cores 4
Dis	4 core	The number of threads 8
Vic	2.2ghz	Floating Point Unit Integrated
Mem	32nm	Level 1 cache size ? 4 x 32 KB instruction caches 4 x 32 KB data caches
Mem	5 GT/s	Level 2 cache size ? 4 x 256 KB
Mem	Socket	Level 3 cache size 8 MB shared cache
Op	Caches	Cache latency 4 (L1 cache) 11 (L2 cache) 25 (L3 cache)
So	L1 dat	Physical memory 32 GB
Int	L1 inst	Multiprocessing Uniprocessor
Co	L2: 10	<ul style="list-style-type: none"> ◦ MMX instructions ◦ SSE / Streaming SIMD Extensions ◦ SSE2 / Streaming SIMD Extensions 2 ◦ SSE3 / Streaming SIMD Extensions 3 ◦ SSSE3 / Supplemental Streaming SIMD Extensions 3 ◦ SSE4 / SSE4.1 + SSE4.2 / Streaming SIMD Extensions 4 ? ◦ AES / Advanced Encryption Standard instructions ◦ AVX / Advanced Vector Extensions ◦ EM64T / Extended Memory 64 technology / Intel 64 ? ◦ NX / XD / Execute disable bit ? ◦ HT / Hyper-Threading technology ? ◦ TBT 2.0 / Turbo Boost technology 2.0 ? ◦ VT-x / Virtualization technology ? ◦ VT-d / Virtualization for directed I/O
	Bus sp	

Processor / Chipset

Architecture / Microarchitecture

Processor	Microarchitecture	Sandy Bridge
	Processor core	Sandy Bridge
	Core stepping	D2 (Q1HK, SR00B)

CPU-Z

Processor

- Caches | Mainboard | Memory | SPD | Graphics | About

Processor

- Name: Intel Core i5 2500K
- Processor Name: Sandy Bridge
- Max TDP: 95 W
- Package: Socket 1155 LGA
- Technology: 32 nm
- Core Voltage: 1.304 V
- Classification: Intel(R) Core(TM) i5-2500K CPU @ 3.30GHz
- Family: 6
- Model: A
- Stepping: 7
- Family: 6
- Ext. Model: 2A
- Revision: D2
- Instructions: MMX, SSE (1, 2, 3, 3S, 4.1, 4.2), EM64T, VT-x, AES, AVX

Cache (Core #0)

Speed: 4327.3 MHz	Cache:	
Multplier: x 42.0	L1 Data: 4 x 32 KBytes	8-way
Speed: 103.0 MHz	L1 Inst.: 4 x 32 KBytes	8-way
id FSB:	Level 2: 4 x 256 KBytes	8-way
ection:	Level 3: 6 MBytes	12-way

Processor #1 Cores: 4 Threads: 4

CPU-Z Version 1.58 Validate OK

CPU-Z

General

Type: DDR3	Channels #: Dual
Size: 8192 MBytes	DC Mode:
NB Frequency:	

Timings

DRAM Frequency: 961.7 MHz
FSB:DRAM: 1:7
CAS# Latency (CL): 9.0 clocks
RAS# to CAS# Delay (tRCD): 11 clocks
RAS# Precharge (tRP): 9 clocks
Cycle Time (tRAS): 27 clocks
Bank Cycle Time (tRC):
Command Rate (CR): 2T
DRAM Idle Timer:
Total CAS# (tRDRAM):
Row To Column (tRCD):

CPU-Z Version 1.58 Validate OK

CPU-Z

Memory Slot Selection

Slot #2: DDR3
Module Size: 4096 MB
Max Bandwidth: PC3-10700 (1066MHz)
Manufacturer: Team Group
Part Number: Xtreem-LV
Serial Number:

Timings Table

JEDEC #2
Frequency: 533 MHz
CAS# Latency: 7.0
RAS# to CAS#: 7
RAS# Precharge: 7
tRAS: 20
tRC: 27
Command Rate:
Voltage: 1.50 V

CPU-Z Version 1.58 Validate OK

- Supplemental Streaming SIMD Extensions 4
- SSE4 / SSE4.1 + SSE4.2 / Streaming SIMD Extensions 4
 - AES / Advanced Encryption Standard instructions
 - AVX / Advanced Vector Extensions
 - EM64T / Extended Memory 64 technology / Intel 64
 - NX / XD / Execute disable bit
 - HT / Hyper-Threading technology
 - TBT 2.0 / Turbo Boost technology 2.0
 - VT-x / Virtualization technology
 - VT-d / Virtualization for directed I/O

Memory Hierarchy Levels



*Capacity
Access Time
Cost*

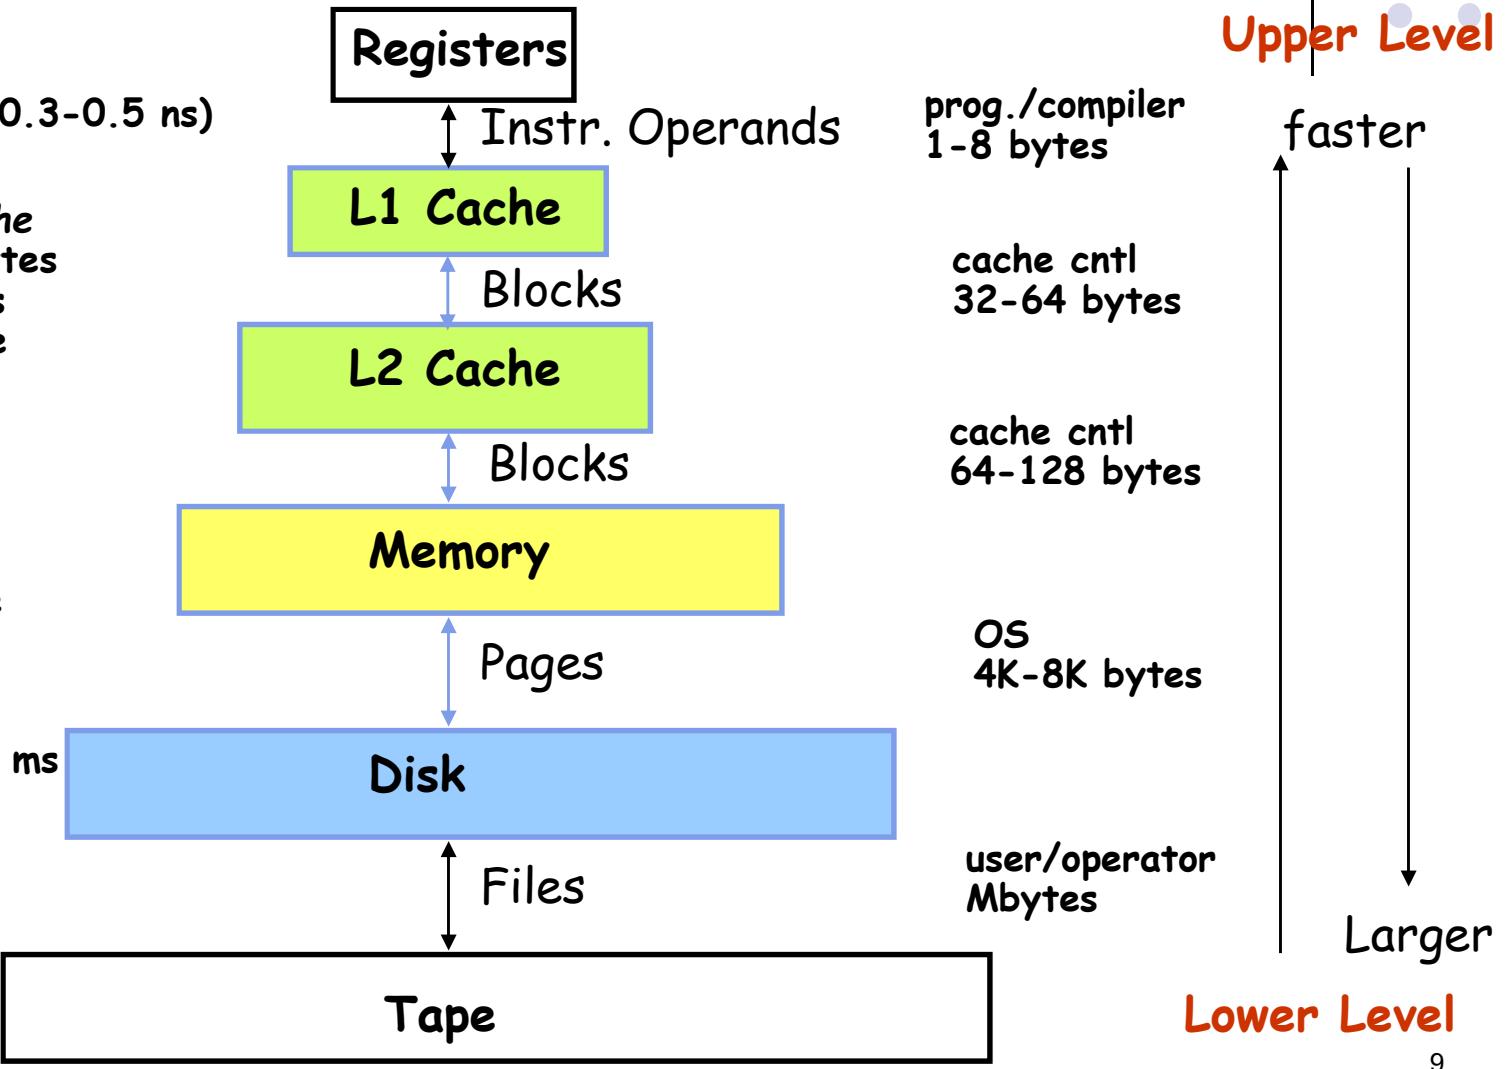
*CPU Registers
100s Bytes
300 - 500 ps (0.3-0.5 ns)*

*L1 and L2 Cache
10s-100s K Bytes
~1 ns - ~10 ns
\$1000s/ GByte*

*Main Memory
G Bytes
80ns- 200ns
~ \\$100/ GByte*

*Disk
10s T Bytes, 10 ms
(10,000,000 ns)
~ \\$1 / GByte*

*Tape
infinite
sec-min
~\\$1 / GByte*



Principle of Locality



- Programs access a small proportion of their address space at any time
- **Temporal locality**
 - Items accessed recently are likely to be accessed again soon
 - e.g., instructions in a loop, induction variables
- **Spatial locality**
 - Items near those accessed recently are likely to be accessed soon
 - E.g., sequential instruction access, array data

Control structure



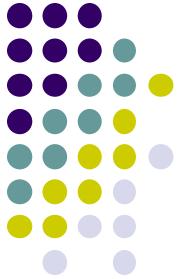
- *Loop*

```
cout << "Adding the numbers\n";
for (i = 0; i < v1.size(); i++) {
    v3[i] = v1[i] + v2[i];
}
```

- *Sequential*

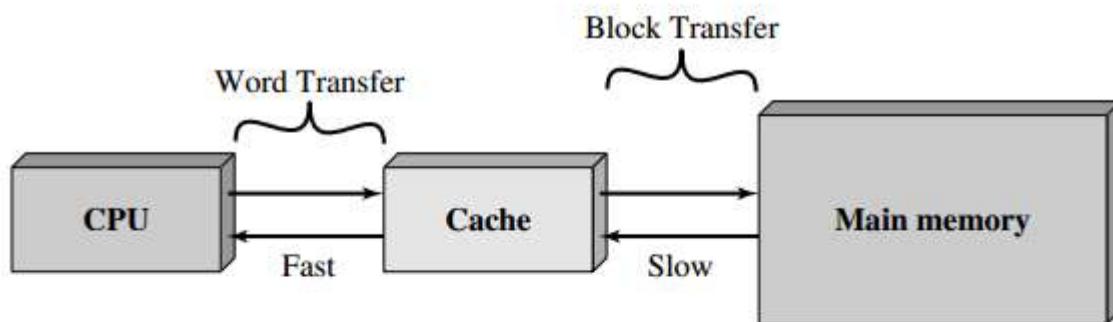
```
lw $t2, 0($t3)
lw $t3, 4($t3)
beq $t2, $t3, Label #assume not equal
add $t5, $t2, $t3
sw $t5, 8($t3)
Label:      j 32
```

Taking Advantage of Locality

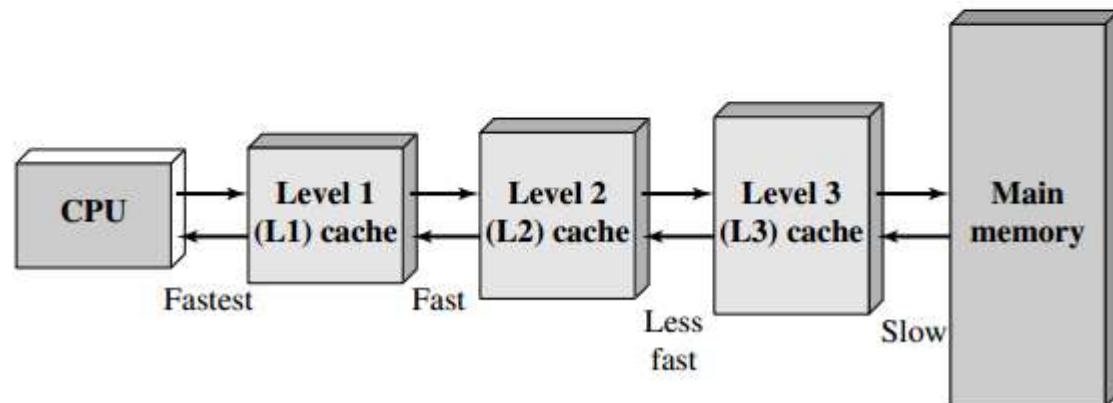


- Memory hierarchy
- Store everything on disk
- Copy recently accessed (and nearby) items from disk to smaller DRAM memory
 - Main memory
- Copy more recently accessed (and nearby) items from DRAM to smaller SRAM memory
 - Cache memory attached to CPU

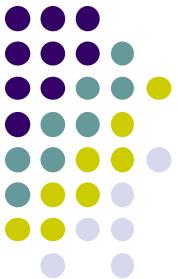
Cache position



(a) Single cache

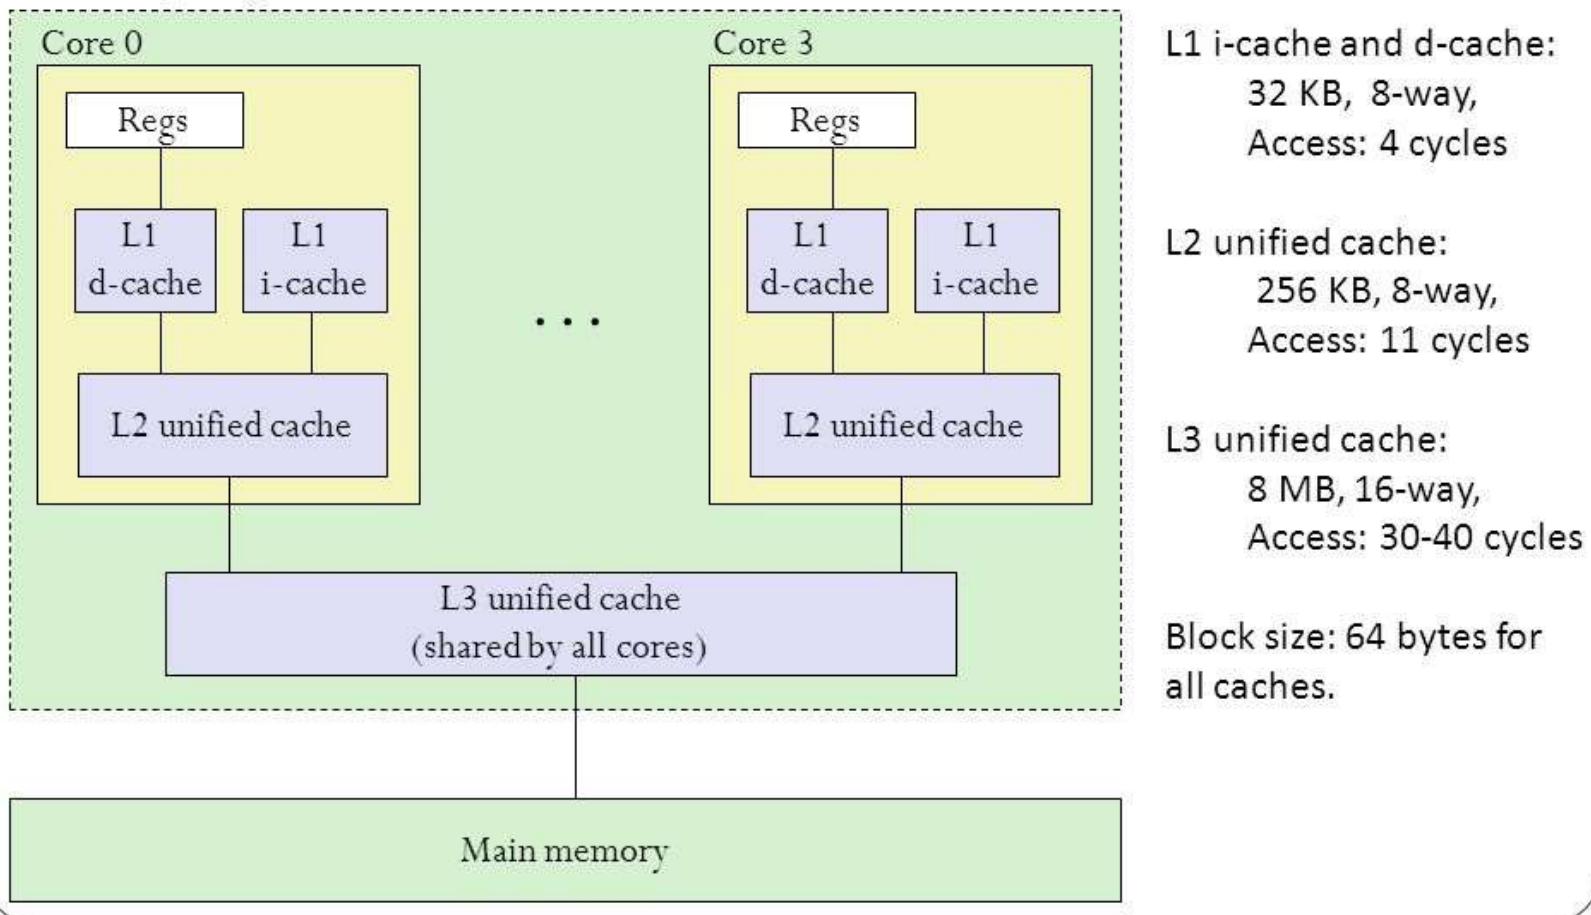


(b) Three-level cache organization

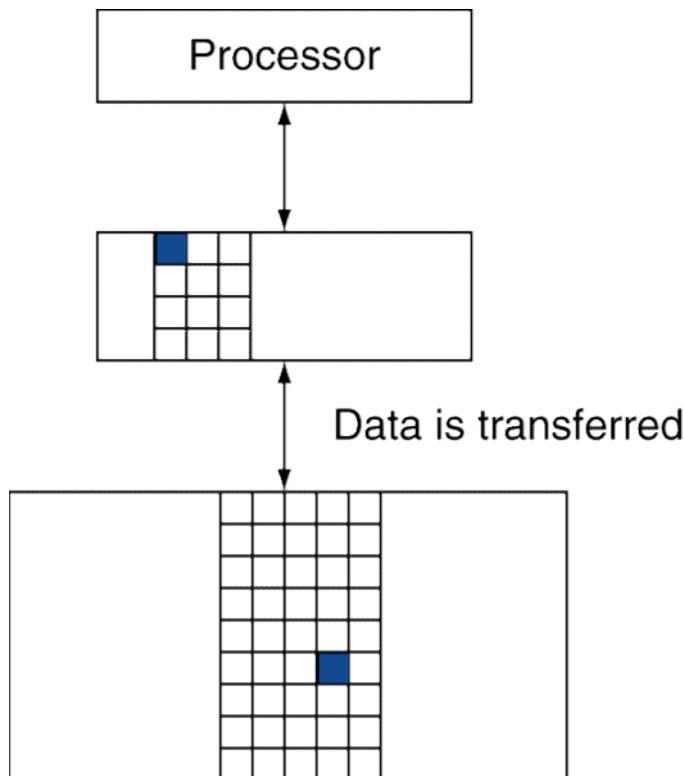


Intel Core i7 Cache Hierarchy

Processor package

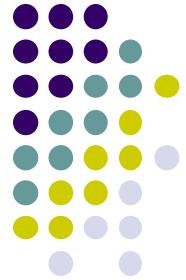


Memory Hierarchy Levels

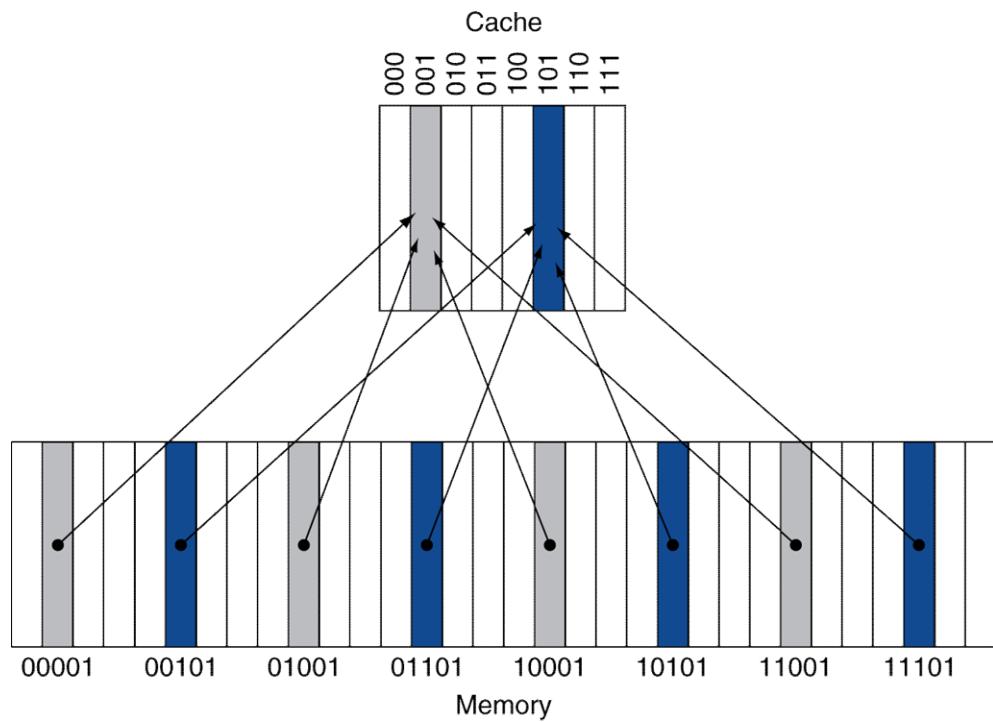


- Block (aka line): unit of copying
 - May be multiple words
- If accessed data is present in upper level
 - Hit: access satisfied by upper level
 - Hit ratio: hits/accesses
- If accessed data is absent
 - Miss: block copied from lower level
 - Time taken: miss penalty
 - Miss ratio: misses/accesses
 $= 1 - \text{hit ratio}$
 - Then accessed data supplied from upper level

Direct Mapped Cache



- Location determined by address
- Direct mapped: only one choice
 - (Block address) modulo (#Blocks in cache)

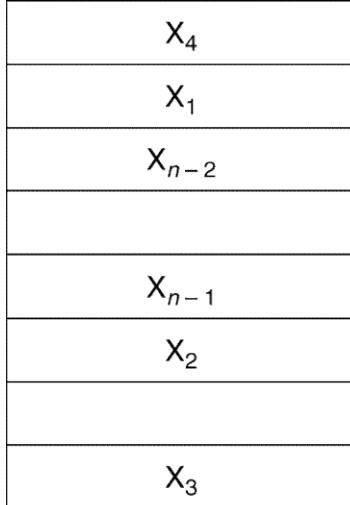


- #Blocks is a power of 2
- Use low-order address bits

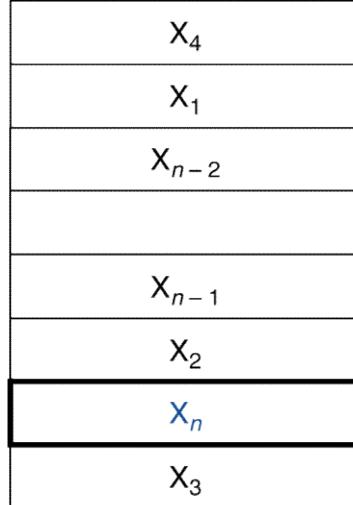


Cache Memory

- Cache memory
 - The level of the memory hierarchy closest to the CPU
- Given accesses X_1, \dots, X_{n-1}, X_n



a. Before the reference to X_n



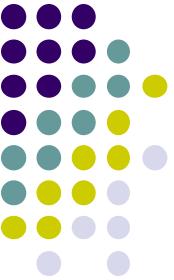
b. After the reference to X_n

- How do we know if the data is present?
- Where do we look?

Tags and Valid Bits



- How do we know which particular block is stored in a cache location?
 - Store block address as well as the data
 - Actually, only need the high-order bits
 - Called the tag
- What if there is no data in a location?
 - Valid bit: 1 = present, 0 = not present
 - Initially 0



Cache Example

- 8-blocks, 1 word/block, direct mapped
- Initial state

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	N		
111	N		

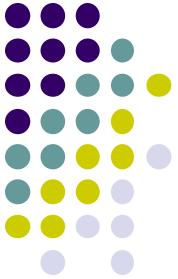


Cache Example (cont'd)

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Miss	110

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Cache Example (cont'd)



Word addr	Binary addr	Hit/miss	Cache block
26	11 010	Miss	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example (cont'd)

Word addr	Binary addr	Hit/miss	Cache block
22	10 110	Hit	110
26	11 010	Hit	010

Index	V	Tag	Data
000	N		
001	N		
010	Y	11	Mem[11010]
011	N		
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example (cont'd)

Word addr	Binary addr	Hit/miss	Cache block
16	10 000	Miss	000
3	00 011	Miss	011
16	10 000	Hit	000

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	11	Mem[11010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		



Cache Example (cont'd)

Word addr	Binary addr	Hit/miss	Cache block
18	10 010	Miss	010

Index	V	Tag	Data
000	Y	10	Mem[10000]
001	N		
010	Y	10	Mem[10010]
011	Y	00	Mem[00011]
100	N		
101	N		
110	Y	10	Mem[10110]
111	N		

Mapping from memory to cache



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	...
0	1	2	3	0	1	2	3	0	1	2	3	0	1	2	3	...
Block offset				Block offset				Block offset				Block offset				...
0				1				2				3				...

Program is divided into equal chunks of the same size as the cache block

`BlockIndex = x / block_size`

`BlockOffset = x mod block_size`

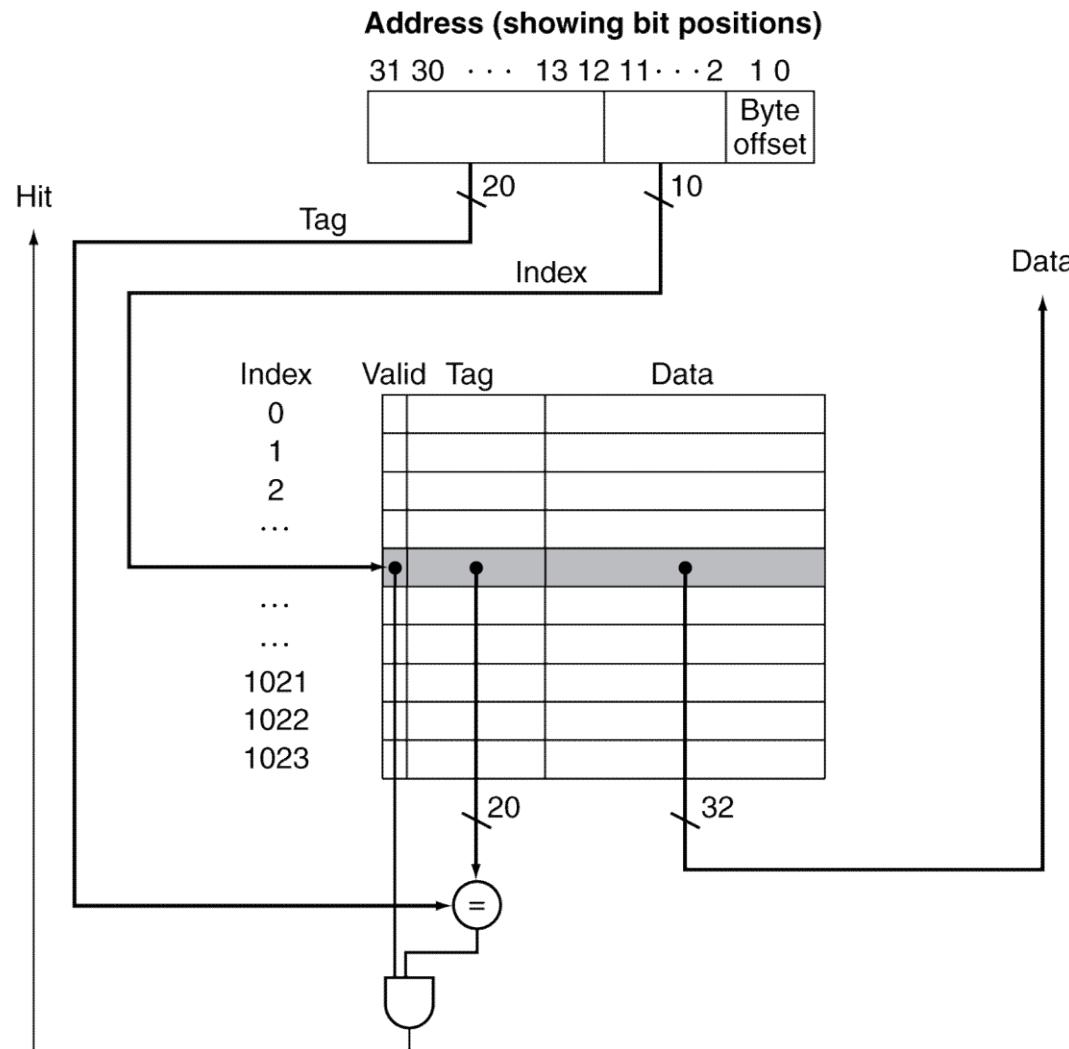
For example: $x=14$

- $\text{BlockIndex} = 14 / 4 = 3$
- $\text{BlockOffset} = 14 \bmod 4 = 2$

Map: use `BlockIndex` for calculating the Block in the cache



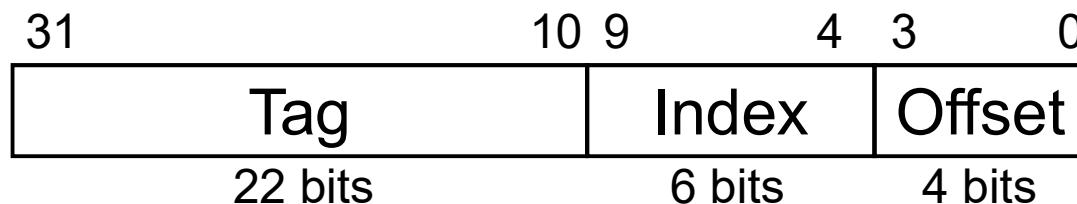
Address Subdivision





Example: Larger Block Size

- 64 blocks, 16 bytes/block, 32 bit address regs
 - To what block number does address 1200 map?
- Offset=?
- Index=?
- Tag=?





Cache usage protocol

- Identify the index, tag, offset of the reference
- Check the valid bit
 - If $v=valid$, check the tag
 - If two tags are equal => return the data (cache hit)
 - Else raise a cache miss => Load the data into the WHOLE slot, update tag and $v=valid$, return the data
 - If $v=invalid$, raise a cache miss
 - Load the data into the WHOLE slot, update tag and $v=valid$, return the data

Example



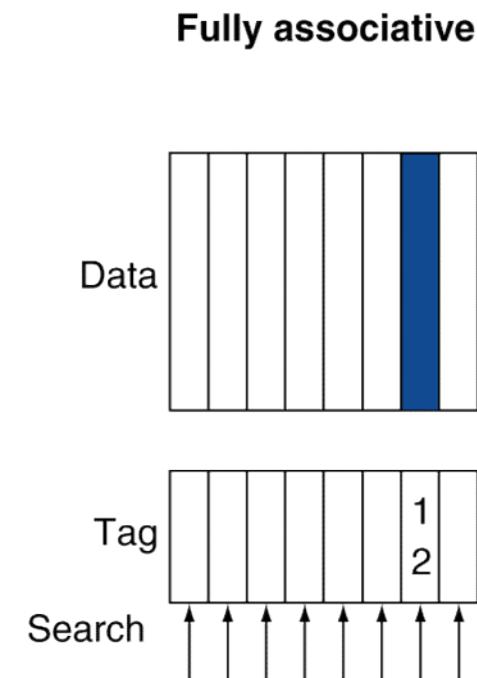
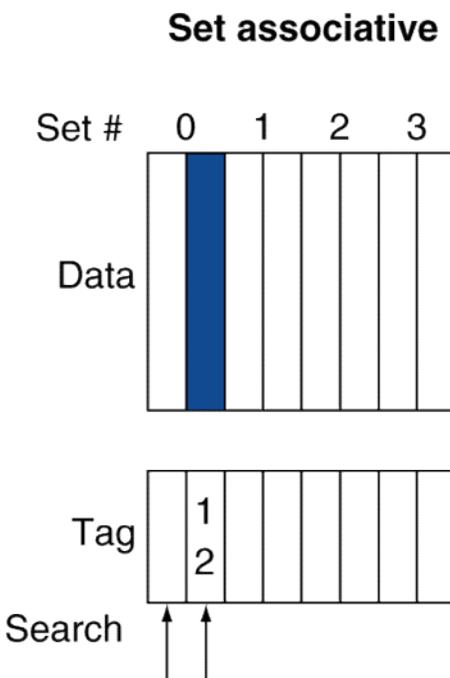
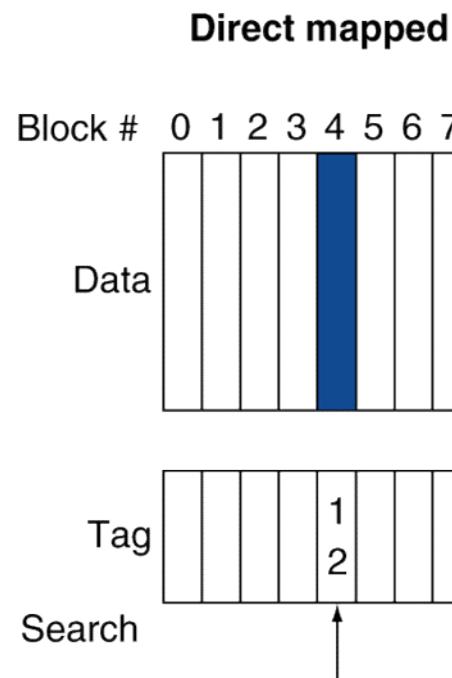
- Suppose a cache 16KB, slot size=32 bytes, address reg=32 bit. Identify the bits for tag, index and offset (ignore tag and valid bit).
- Given the above cache with direct map
 - Cache with 64 slots, slot size=32 bytes, address reg=16 bit
 - Show hit/miss for the reference sequence of MIPS instructions: 184, 188, 192, 196, 200, 204, 208, 212, 216, 192, 196, 200, 204.

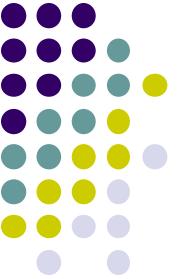


Associative Caches

- Fully associative
 - Allow a given block to go in **any cache** entry
 - Requires all entries to be searched at once
 - Comparator per entry (expensive)
- n -way set associative
 - Each set contains n entries
 - Block number determines which set
 - (Block number) modulo (#Sets in cache)
 - Search all entries in a given set at once
 - n comparators (less expensive)

Associative Cache Example





Spectrum of Associativity

- For a cache with 8 entries

**One-way set associative
(direct mapped)**

Block	Tag	Data
0		
1		
2		
3		
4		
5		
6		
7		

Two-way set associative

Set	Tag	Data	Tag	Data
0				
1				
2				
3				

Four-way set associative

Set	Tag	Data	Tag	Data	Tag	Data	Tag	Data
0								
1								

Eight-way set associative (fully associative)

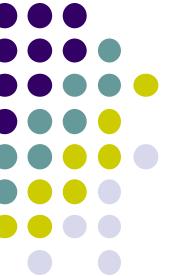
Tag	Data														



Associativity Example

- Compare 4-block caches
 - Direct mapped, 2-way set associative, fully associative
 - Block access sequence: 0, 8, 0, 6, 8
- Direct mapped

Block address	Cache index	Hit/miss	Cache content after access			
			0	1	2	3
0	0	miss	Mem[0]			
8	0	miss	Mem[8]			
0	0	miss	Mem[0]			
6	2	miss	Mem[0]		Mem[6]	
8	0	miss	Mem[8]		Mem[6]	



Associativity Example

- 2-way set associative

Block address	Cache index	Hit/miss	Cache content after access	
			Set 0	Set 1
0	0	miss	Mem[0]	
8	0	miss	Mem[0]	Mem[8]
0	0	hit	Mem[0]	Mem[8]
6	0	miss	Mem[0]	Mem[6]
8	0	miss	Mem[8]	Mem[6]

- Fully associative

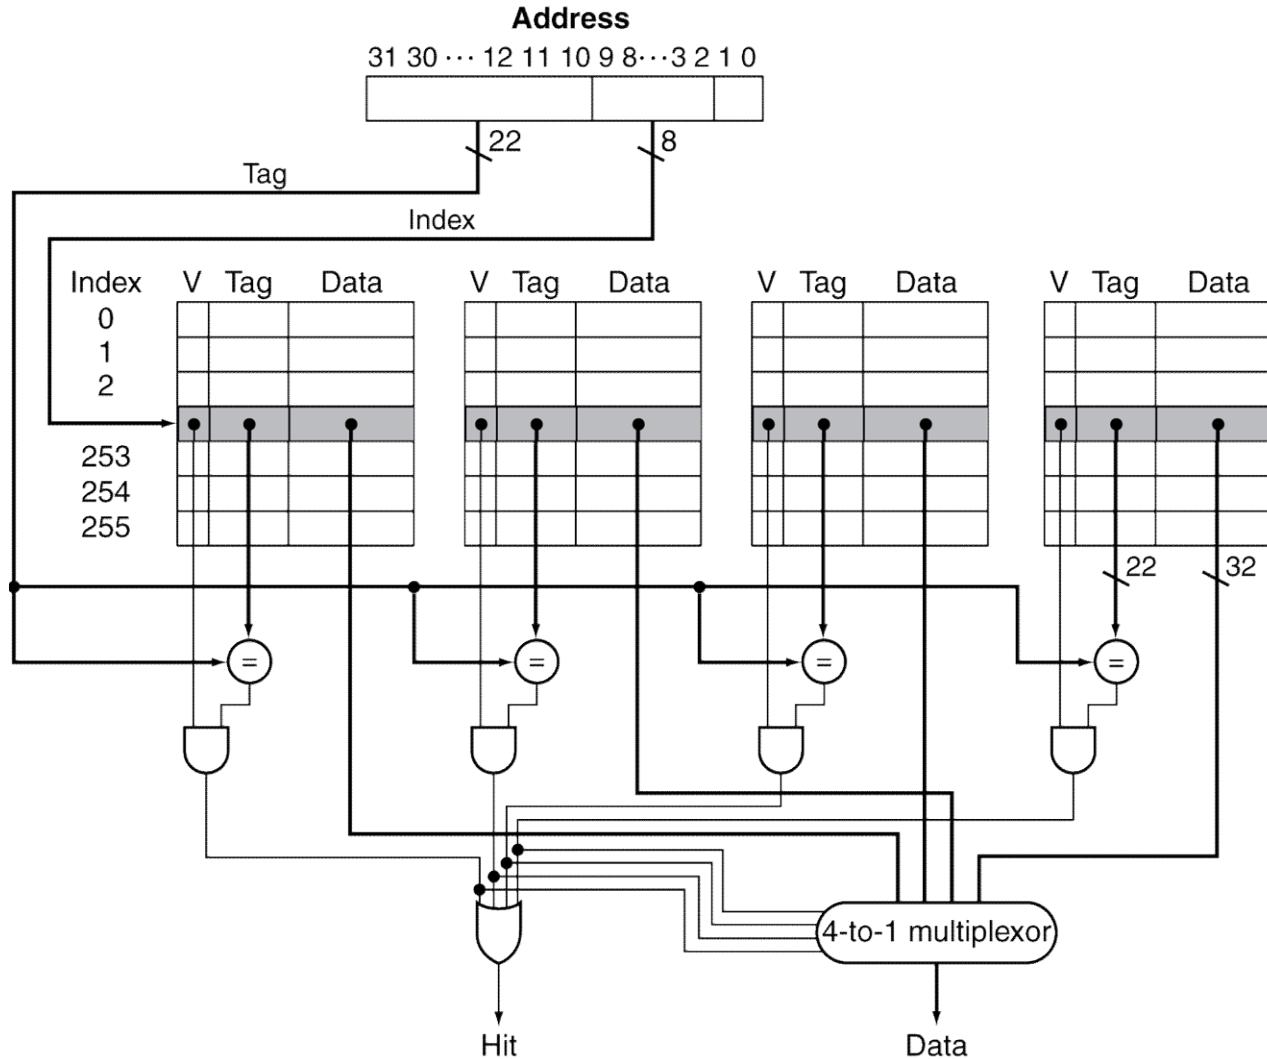
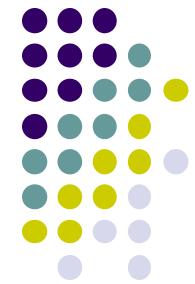
Block address		Hit/miss	Cache content after access		
			Mem[0]	Mem[8]	Mem[6]
0		miss	Mem[0]		
8		miss	Mem[0]	Mem[8]	
0		hit	Mem[0]	Mem[8]	
6		miss	Mem[0]	Mem[8]	Mem[6]
8		hit	Mem[0]	Mem[8]	Mem[6]



How Much Associativity

- Increased associativity decreases miss rate
 - But with diminishing returns
- Simulation of a system with 64KB D-cache, 16-word blocks, SPEC2000
 - 1-way: 10.3%
 - 2-way: 8.6%
 - 4-way: 8.3%
 - 8-way: 8.1%

Set Associative Cache Organization





Exercise

- Given a 4-way cache size =128KB (for data), block size=32 bytes, the address register=32 bit.
 - Identify the components of the register.
 - Given a reference x=1864, identify the components



Replacement Policy

- Direct mapped: no choice
- Set associative
 - Prefer non-valid entry, if there is one
 - Otherwise, choose among entries in the set
- Least-recently used (LRU)
 - Choose the one unused for the longest time
 - Simple for 2-way, manageable for 4-way, too hard beyond that
- Random
 - Gives approximately the same performance as LRU for high associativity



Block Size Considerations

- Larger blocks should reduce miss rate
 - Due to spatial locality
- But in a fixed-sized cache
 - Larger blocks \Rightarrow fewer of them
 - More competition \Rightarrow increased miss rate
 - Larger blocks \Rightarrow pollution
- Larger miss penalty
 - Can override benefit of reduced miss rate
 - Early restart and critical-word-first can help



Cache Misses

- On cache hit, CPU proceeds normally
- On cache miss
 - Stall the CPU pipeline
 - **Fetch** block from next level of hierarchy
 - Instruction cache miss
 - Restart instruction fetch
 - Data cache miss
 - Complete data access



Write-Through

- On data-write hit, could just update the block in cache
 - But then cache and memory would be inconsistent
- Write through: also update memory
- But makes writes take longer
 - e.g., if base CPI = 1, 10% of instructions are stores, write to memory takes 100 cycles
 - Effective CPI = $1 + 0.1 \times 100 = 11$
- Solution: write buffer
 - Holds data waiting to be written to memory
 - CPU continues immediately
 - Only stalls on write if write buffer is already full



Write-Back

- Alternative: On data-write hit, just update the block in cache
 - Keep track of whether each block is dirty
- When a dirty block is replaced
 - Write it back to memory
 - Can use a write buffer to allow replacing block to be read first



Write Allocation

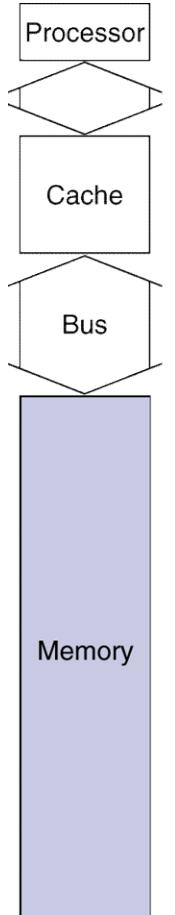
- What should happen on a write miss?
- Alternatives for write-through
 - Allocate on miss: fetch the block
 - Write around: don't fetch the block
 - Since programs often write a whole block before reading it (e.g., initialization)
- For write-back
 - Usually fetch the block



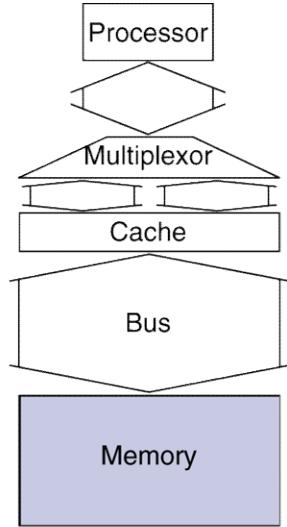
Main Memory Supporting Caches

- Use DRAMs for main memory
 - Fixed width (e.g., 1 word)
 - Connected by fixed-width clocked bus
 - Bus clock is typically slower than CPU clock
- Example cache block read
 - 1 bus cycle for address transfer
 - 15 bus cycles per DRAM access
 - 1 bus cycle per data transfer
- For 4-word block, 1-word-wide DRAM
 - Miss penalty = $1 + 4 \times 15 + 4 \times 1 = 65$ bus cycles
 - Bandwidth = 16 bytes / 65 cycles = 0.25 B/cycle

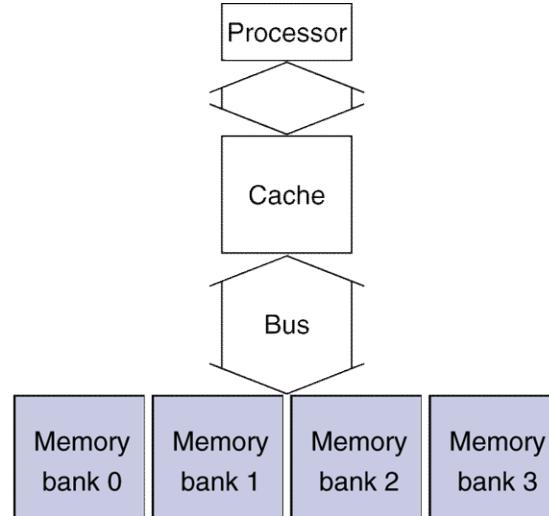
Increasing Memory Bandwidth



a. One-word-wide
memory organization



b. Wider memory organization



c. Interleaved memory organization

- 4-word wide memory
 - Miss penalty = $1 + 15 + 1 = 17$ bus cycles
 - Bandwidth = $16 \text{ bytes} / 17 \text{ cycles} = 0.94 \text{ B/cycle}$
- 4-bank interleaved memory
 - Miss penalty = $1 + 15 + 4 \times 1 = 20$ bus cycles
 - Bandwidth = $16 \text{ bytes} / 20 \text{ cycles} = 0.8 \text{ B/cycle}$



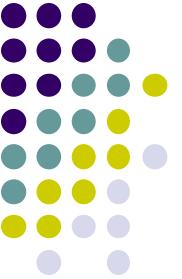
Average Access Time

- Hit time is also important for performance
- Average memory access time (AMAT)
 - $AMAT = \text{Hit time} + \text{Miss rate} \times \text{Miss penalty}$
- Effective CPI
 - $ECPI = CPI(\text{ideal}) + \text{miss rate}_1 \times \text{Miss penalty} + \dots$
- Example
 - CPU with 1ns clock, hit time = 1 cycle, miss penalty = 20 cycles, I-cache miss rate = 5%
 - $AMAT = 1 + 0.05 \times 20 = 2$
 - 2 cycles per instruction

Cache Performance Example



- Given
 - I-cache (instruction) miss rate = 2%
 - D-cache (data) miss rate = 4%
 - Miss penalty = 100 cycles
 - Base CPI (ideal cache) = 2
 - Load & stores are 36% of instructions
- Miss cycles per instruction
 - I-cache: $0.02 \times 100 = 2$
 - D-cache: $0.36 \times 0.04 \times 100 = 1.44$
- Actual CPI = $2 + 2 + 1.44 = 5.44$
 - Ideal CPU is $5.44/2 = 2.72$ times faster



Multilevel Caches

- Primary cache attached to CPU
 - Small, but fast
- Level-2 cache services misses from primary cache
 - Larger, slower, but still faster than main memory
- Main memory services L-2 cache misses
- Some high-end systems include L-3 cache



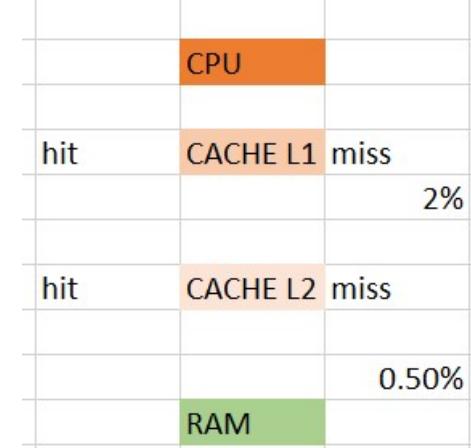
Multilevel Cache Example

- Given
 - CPU base CPI = 1, clock rate = 4GHz
 - Miss rate/instruction = 2%
 - Main memory access time = 100ns
- With just primary cache
 - CPU clock cycle= $1/(4*10^9) = \frac{1}{4}$ ns=0.25ns
 - Miss penalty = $100\text{ns}/0.25\text{ns} = 400$ cycles
 - Effective CPI = $1 + 0.02 \times 400 = 9$



Example (cont.)

- Now add L-2 cache
 - Access time = 5ns
 - Global miss rate to main memory = 0.5%
- Primary miss with L-2 hit
 - Penalty = $5\text{ns}/0.25\text{ns} = 20$ cycles
- Primary miss with L-2 miss
 - Extra penalty = 500 cycles
- CPI = $1 + 0.02 \times 20 + 0.005 \times 500 = 3.9$
- Performance ratio = $9/3.9 = 2.3$





Performance Summary

- When CPU performance increased
 - Miss penalty becomes more significant
- Decreasing base CPI
 - Greater proportion of time spent on memory stalls
- Increasing clock rate
 - Memory stalls account for more CPU cycles
- Can't neglect cache behavior when evaluating system performance